The Capture Problem: How Infrastructure Departments Became the Biggest Barrier to Enterprise AI ROI

Pat Brenner

March 2026

Abstract

The biggest barrier to enterprise AI ROI is not model capability, data readiness, or regulation. It is the organizational capture of AI by infrastructure departments whose institutional incentives prioritize risk minimization over value creation. This paper argues that a governance pattern well-documented in the history of enterprise technology—the seizure of strategic ownership by departments that control infrastructure—is repeating with AI, and that the consequences this time are more severe because AI permeates internal business processes, not just customer-facing channels.

Drawing on the parallel case of web ownership struggles in the late 1990s and early 2000s, the paper traces how IT departments leveraged infrastructure control to claim authority over strategy, content, and functionality—domains where they had the least proximity to customers and business outcomes. It identifies the same structural dynamic in today's enterprise AI landscape, compounded by the broader surface area AI touches, a fundamental velocity mismatch between AI-native development cycles and traditional IT governance timelines, and misaligned incentive systems that reward risk avoidance over value delivery.

The paper introduces a four-form governance taxonomy—strategic, process, compliance, and value governance—and argues that the current enterprise AI model concentrates all four under IT ownership, despite IT holding strong competence in only one or two. Empirical data from MIT, IBM, Deloitte, McKinsey, and the Cloud Security Alliance consistently shows that governance maturity is the strongest predictor of AI readiness and ROI, but that the dominant form of governance—centralized, IT-controlled, process-heavy—is itself the primary bottleneck. Case studies of DBS Bank and IBM illustrate what business-owned AI governance produces when properly structured.

The paper proposes a governance model in which business units own strategic and value governance, IT provides process governance for technical execution, and Responsible AI owns compliance governance—a clear separation of concerns in which each function governs its domain with full cost transparency across all stakeholders. This is the same resolution that ultimately emerged from the web wars, but which organizations cannot afford to reach through slow, organic correction. The competitive costs of misaligned AI governance compound with every quarter, and the gap between organizations that have resolved the ownership question and those that have not is widening, not narrowing.

Introduction: We've Seen This Movie Before

Enterprise AI is underperforming expectations. Despite unprecedented investment in models, platforms, talent, and professional services, the anticipated productivity gains

and cost reductions remain elusive for most organizations. The standard explanations are familiar: the technology is immature, data isn't ready, regulation creates uncertainty, change management is hard. Each of these contains some truth. None of them is the primary problem.

The primary problem is organizational. Specifically, it is the systematic capture of AI strategy, prioritization, vendor selection, and deployment by infrastructure-oriented technology departments—CIO organizations, enterprise IT, and their associated consulting and professional services ecosystems. This capture is not the result of malice or incompetence. It is the predictable consequence of an institutional dynamic that has played out before, with well-documented costs and a well-documented resolution.

The last time this happened, the technology was the World Wide Web.

Before proceeding, a note on terminology. This paper uses the word "governance" extensively, and it is important to be precise about what the term means—because much of the dysfunction this paper diagnoses stems from organizations treating governance as a single activity when it is, in fact, several distinct disciplines operating under a shared umbrella. Governance, as used here, refers to the enterprise-wide operational discipline of making AI decisions, executing them, and measuring whether they worked. Within that umbrella, four distinct forms of governance are at play, and distinguishing them is essential to understanding why the current model fails and what a better model requires.

Strategic governance determines who owns AI strategy, prioritization, and decision-making authority—which initiatives to pursue, in what order, and measured against which outcomes. This is the governance form most central to this paper's argument. The capture problem is fundamentally a strategic governance failure: the function that controls AI prioritization lacks proximity to the business outcomes that AI is meant to produce.

Process governance encompasses the formal review cycles, approval gates, and operational workflows through which AI initiatives move from concept to deployment. Architecture reviews, security assessments, staged rollouts, and resource allocation decisions all fall within process governance. This is the governance form most often cited when organizations complain that "governance is too slow"—but the slowness is typically a consequence of who controls the process, not an inherent property of the process itself.

Compliance governance is the function of ensuring AI systems meet legal, regulatory, ethical, and reputational standards. This is the domain of Responsible AI teams, legal departments, and regulatory affairs. It is a necessary and legitimate function—but it is an assessment and advisory function, not a strategic or prioritization

function. Problems arise when compliance governance expands beyond its core mandate to become the de facto decision-making authority for all AI activity.

Value governance is the ongoing measurement of whether AI deployments deliver expected business outcomes at acceptable cost. It is the Control phase of the continuous improvement cycle: comparing actual results against projections, with full transparency into all cost components—including IT development costs and ongoing maintenance—visible to all stakeholders. Value governance is the form most conspicuously absent from current enterprise AI programs. Organizations measure deployment velocity (how fast), compliance status (how safe), and technical capability (how sophisticated)—but rarely measure business value delivery with the rigor and transparency that any other capital investment would demand.

The central argument of this paper is that the current enterprise AI governance model concentrates all four forms under IT ownership. This concentration fails because IT holds strong competence in process governance and contributes meaningfully to compliance governance, but lacks the business proximity required for effective strategic governance and the outcome accountability required for effective value governance. The proposed alternative does not diminish IT's role—it clarifies it. Business functions own strategic and value governance. IT provides process governance for technical execution. Responsible AI owns compliance governance. Each function governs its domain, with full transparency across all stakeholders. The result is not less governance but better-structured governance—an enterprise-wide operational discipline in which every AI investment is prioritized by value, executed by the appropriate technical function, assessed for risk, and measured against outcomes.

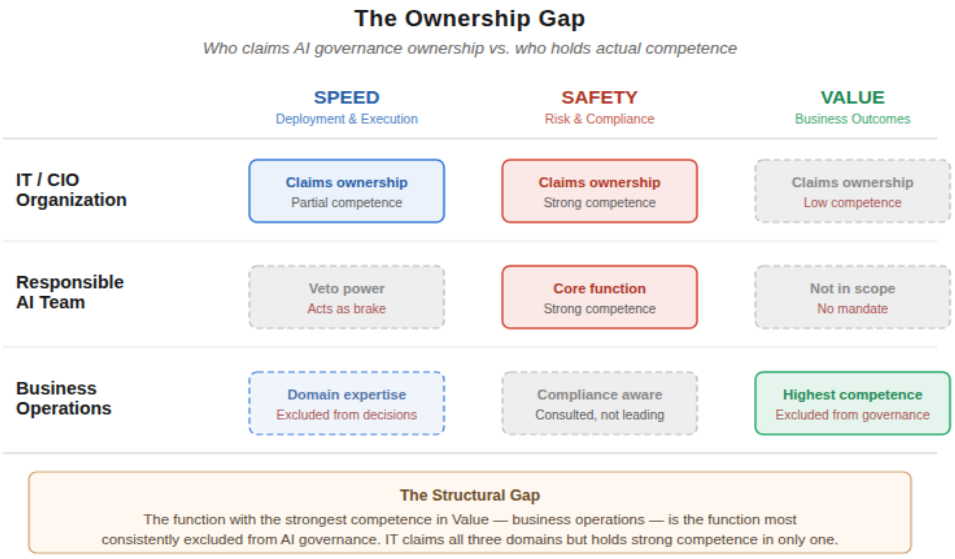

Figure 1. The Ownership Gap — Who claims AI governance ownership vs. who holds actual competence.

The Web Wars: A Case Study in Infrastructure Capture

The Window of Business-Led Innovation

When the commercial web emerged in the mid-1990s, there was no established playbook for organizational ownership. The web was new, strange, and its potential was visible primarily to people who worked at the boundary between organizations and their audiences: marketing teams, product groups, publishing operations, and communications departments. These teams moved first because they could see the possibilities most clearly. They built early websites, experimented with interactive content, tested what resonated with users, and iterated rapidly in an environment where failure was cheap and learning was fast.

This early period was characterized by direct contact between business objectives and web execution. The people building web experiences understood their audiences because serving those audiences was their primary job. Decisions about content, functionality, and user experience were made by people with domain expertise and customer proximity. The results were uneven—some experiments succeeded, many failed—but the feedback loops were tight and the learning was genuine.

The Infrastructure Claim

As web presence grew in scale and organizational importance, a predictable thing happened: IT departments asserted control. The logic was reasonable on its face—web infrastructure required servers, security, networking, and integration with enterprise systems, all of which fell within IT's existing domain. IT could credibly argue that it needed to govern the technical infrastructure of the web.

But infrastructure governance expanded, as it tends to do, into areas far beyond infrastructure. IT departments began controlling timelines, dictating functionality, making content decisions, choosing vendors, and setting strategic priorities for web presence. The mechanism was straightforward: if you control the infrastructure, you control the deployment pipeline. If you control the deployment pipeline, you have effective veto power over what gets built, when, and how. Infrastructure authority became strategic authority not because IT had superior strategic judgment, but because infrastructure is upstream of everything else.

The result was a prolonged period in which enterprise web strategy was driven by departments optimizing for infrastructure concerns—stability, standardization, risk reduction, and control—rather than business outcomes like customer engagement, conversion, and revenue. The people closest to the customer were subordinated to the people closest to the server rack.

The Resolution and Its Cost

This ownership struggle eventually resolved. The organizations that performed best were those that established clear models in which business units owned web strategy and outcomes while IT provided enabling infrastructure. Marketing owned the brand experience. Product owned the functionality. IT made it all run reliably. This model didn't emerge easily or quickly—it took years of organizational friction, lost competitive advantage, and missed market windows. But it emerged because the alternative—infrastructure departments driving strategy for channels they didn't deeply understand—consistently destroyed value.

The companies that recognized this pattern earliest gained years of competitive advantage. The companies that let it play out organically paid for the education with lost time, lost customers, and lost revenue.

The AI Replay: Same Pattern, Higher Stakes

The New Capture

The same structural dynamic is now playing out with enterprise AI, accelerated and amplified. CIO organizations and enterprise IT departments are claiming ownership of AI strategy, prioritization, vendor selection, and deployment across the enterprise. As with the web, the initial claim is grounded in legitimate infrastructure concerns: AI requires compute, data pipelines, model hosting, security, and integration. These are real requirements that fall within technology departments' competence.

And as with the web, infrastructure authority is expanding into strategic territory. IT-led AI teams are choosing which business processes to automate and in what order. They are selecting vendors based on technical architecture preferences rather than business outcome fitness. They are authoring enterprise AI strategies with limited input from the business functions those strategies are meant to serve. They are defining what "responsible AI" means in ways that prioritize institutional risk minimization over business value creation. In many organizations, AI teams embedded within IT pillars are not merely enabling business process changes—they are driving them directly, making decisions about how business operations should run based on what the technology makes convenient rather than what the business requires.

Why the Stakes Are Higher

The web ownership struggle was painful but bounded. The web was primarily a customer-facing channel. When IT captured web strategy, the damage was largely confined to marketing effectiveness, customer experience, and competitive positioning in digital channels. These were significant costs, but they were contained.

AI is different in a critical respect: it touches internal business processes at every level. AI-driven automation affects order-to-cash workflows, customer service operations, supply chain management, financial planning, HR operations, and dozens of other internal functions. When infrastructure departments capture AI strategy, the domain creep extends into the operational core of the business. IT teams are not just controlling a customer-facing channel—they are reshaping how the entire organization works, often without sufficient understanding of the processes they are modifying.

This expanded surface area means the cost of misaligned governance is proportionally larger. A bad web strategy lost customers. A bad AI strategy loses operational efficiency, employee productivity, and process integrity across the enterprise.

The Capture in Practice: Composite Patterns from the Field

The dynamics described above are not theoretical. They manifest in recognizable, recurring patterns across industries. The following composites are drawn from widely reported enterprise AI experiences and patterns documented across the research literature. Names and identifying details have been changed, but the structural dynamics are faithful to the originals. Each illustrates a different facet of infrastructure capture and its consequences.

The Vendor Selection Override. A consumer goods company's customer operations team spent six months evaluating AI platforms for its contact center—the largest customer touchpoint in the company. The team ran structured pilots, scored vendors against business outcome criteria, and identified a clear frontrunner that outperformed alternatives on resolution accuracy, agent experience, and cost per interaction. When the recommendation reached the CIO organization for approval, it was rejected. The IT architecture team selected a different vendor—one that scored lower on every business metric but aligned with the enterprise's existing cloud platform and the CIO's stated preference for vendor consolidation. The rationale was infrastructure coherence. The consequence was a deployment that took fourteen months instead of the projected six, required extensive customization to meet basic business requirements the original vendor handled natively, and still underperforms the rejected alternative on the metrics that matter to customers. The business team that understood the problem best was overruled by the team that understood the infrastructure best. The infrastructure is coherent. The business outcome is not.

• • •

The Expertise Inflation Spiral. A large financial services firm launched its enterprise AI program under the CIO's office with a mandate to "go fast." Within eighteen months, the AI team grew from twelve people to over ninety: data scientists, ML engineers, platform architects, a dedicated MLOps team, and a Responsible AI

governance group. The team built an impressive internal AI platform and ran dozens of proofs of concept. When the CFO asked what measurable business value the program had delivered, the answer was difficult to articulate. The proofs of concept demonstrated technical capability, but few had been deployed into production business processes. Those that had were modifications to internal IT workflows—the AI team automating its own adjacent processes rather than the high-value business operations the program was meant to transform. The team lacked the domain expertise to design effective solutions for lending operations, claims processing, or client onboarding without extensive support from business subject matter experts—support that was difficult to secure because the business units had not been consulted on prioritization and did not consider the AI team's chosen projects their most pressing needs. Meanwhile, the program's annual cost had reached eight figures, driven by the specialized talent required to maintain the platform and the professional services engagements needed to compensate for the team's business domain gaps. The AI program was fast, expensive, and largely disconnected from the business outcomes it existed to serve.

• • •

The Governance Bottleneck. A global healthcare company established a Responsible AI board as part of its IT governance structure. The board's mandate was to review all AI deployments for legal, compliance, and reputational risk—a legitimate and necessary compliance governance function. In practice, the board became the single approval gate for every AI initiative in the company, regardless of risk profile. A low-risk internal tool that helped procurement analysts summarize vendor contracts required the same process governance review as a patient-facing diagnostic support system. The review process averaged eleven weeks. For a technology whose development cycles increasingly operate on timescales of days and weeks, eleven weeks of process governance review is not a guardrail—it is a wall. Business units responded predictably: they stopped submitting AI initiatives to the formal process. Shadow AI proliferated, with teams using commercial AI tools outside IT governance entirely. The Responsible AI board's risk-minimization mandate had produced the worst possible outcome—not reduced risk, but invisible risk. The organization lost both the governance visibility it needed and the deployment velocity the business required. The board was effective at one thing: ensuring that the AI initiatives it could see moved slowly. It had no mechanism for ensuring they moved toward value.

• • •

The Process Colonization. A Fortune 500 manufacturer's IT-embedded AI team was tasked with automating the demand forecasting process for its supply chain organization. The supply chain team had specific requirements shaped by years of operational experience: the forecasting model needed to account for promotional lift, regional seasonality patterns, and retailer-specific ordering behaviors that varied by

channel. The AI team, staffed with strong technical talent but limited supply chain experience, built an elegant machine learning solution that optimized for aggregate forecast accuracy. It performed well on the metric the AI team chose to optimize. It performed poorly on the operational reality the supply chain team lived in every day. The model's aggregate accuracy masked significant errors at the SKU-region level that mattered most for production planning and inventory allocation. When the supply chain team raised concerns, the AI team's response was to propose changes to the supply chain planning process to better accommodate the model's outputs. The technology team was not adapting the AI to the business process—it was asking the business process to adapt to the AI. This inversion—where the enabling function dictates operational changes to the business function it is meant to serve—is one of the most corrosive consequences of infrastructure capture. It reverses the proper relationship between technology and the work technology exists to support.

• • •

The Development Velocity Gap. A mid-cap technology company's marketing operations team identified an opportunity to use generative AI to automate the creation of product documentation—a high-volume, repetitive task that consumed significant writer and editor time. Using commercially available AI tools and modern prompt engineering techniques, a two-person team within marketing built a working prototype in three weeks that produced documentation meeting 85% of quality standards on first pass. When the team sought to move the prototype into production, the CIO's office required the project to go through the standard AI development pipeline: architecture review, security assessment, platform integration, QA testing, and staged rollout. The IT-governed process estimated nine months to production deployment. During those nine months, the underlying AI models improved substantially, the prompt engineering techniques the original team had used became outdated, and two competitors launched similar capabilities. The marketing team had demonstrated that AI-native development frameworks could deliver business value in weeks. The infrastructure department's governance processes required months. The technology itself had changed what was possible; the governance model had not changed what was permitted. By the time the approved solution deployed, it was solving yesterday's version of the problem with yesterday's version of the technology.

• • •

These patterns are not exceptional cases. They are the ordinary, everyday consequences of a governance model that places AI strategy ownership in departments optimized for infrastructure management rather than business value creation. The vendor override, the expertise spiral, the governance bottleneck, the process colonization, the velocity gap—these are not five separate problems. They are five symptoms of a single structural condition: infrastructure capture.

The data reinforces what these patterns suggest. Industry surveys consistently find that the majority of organizations struggle to translate AI investments into measurable business outcomes, with some estimates indicating that as few as one in four AI initiatives deliver expected ROI. The gap between AI ambition and governance maturity continues to widen—organizations report scaling AI capabilities while simultaneously acknowledging they lack the governance structures to evaluate whether those capabilities produce value. The pattern is strikingly consistent: deployment activity accelerates while value accountability remains absent. This is not a coincidence. It is the predictable output of a system optimized for deployment velocity and risk avoidance rather than business value.

The Counterfactual: What Business-Owned AI Governance Produces

If infrastructure capture is the diagnosis, the natural question is whether the alternative—business-owned AI strategy with IT as an enabling function—actually produces better outcomes. The evidence, while still emerging, is directional and consistent: organizations that have structured their AI governance around business ownership, cross-functional accountability, and outcome-oriented measurement are outperforming those that have not. Two cases are particularly instructive.

The "Two-in-a-Box" Model: DBS Bank. DBS Bank, Southeast Asia's largest bank, offers perhaps the clearest large-scale example of business-owned AI governance producing measurable results. Rather than centralizing AI under a CIO-led technology function, DBS reorganized around 33 platforms aligned to business segments and products, each jointly led by a business leader and a technology leader sharing the same KPIs. This "two-in-a-box" structure made a structural commitment: technology and business would be co-accountable for outcomes, not siloed into a requester-provider relationship.

The results speak in the language this paper argues matters most—business value. DBS has deployed over 1,500 AI models across more than 370 use cases, and the bank can quantify the economic impact: approximately SGD 1 billion (roughly $768 million) in revenue contribution for 2025, up from SGD 750 million the prior year. These are not theoretical projections or proof-of-concept demonstrations—they are measured outcomes tied to specific business processes: customer engagement models that increased savings uptake by 83%, credit risk models that identified 95% of non-performing loans months before distress, and operational tools that reduced system incidents from technology changes by 81%.

Critically, DBS's model did not sacrifice governance for speed. The bank developed an internal framework called PURE—Purposeful, Unsurprising, Respectful, and Explainable—that governs all data and AI use cases before deployment. But PURE operates as a set of principles embedded within business-led development, not as an

external approval gate controlled by an infrastructure department. Governance is woven into how business teams build, not imposed as a checkpoint they must survive. The bank also reduced model development timelines from 18 months to two or three months—not by bypassing governance, but by making governance a design principle rather than a deployment obstacle.

What makes DBS instructive for this paper's argument is not that it is a technology company leveraging AI. It is a bank—a heavily regulated institution in a conservative industry—that achieved AI-driven business transformation by placing ownership where business outcomes are understood and measured. The joint business-technology leadership model ensured that AI initiatives were selected and prioritized based on business value, that domain expertise shaped solution design, and that accountability for results lived with the people closest to the processes being transformed.

The "Fusion Team" Model: IBM. IBM's internal AI transformation offers a different but structurally compatible illustration. Rather than routing all AI development through centralized IT teams, IBM created what it calls "AI fusion teams"—hybrid groups that combine people with deep business function expertise with technologists from the CIO organization. The design principle is explicit: collapse the handoff chain that traditionally separates business knowledge from technical execution.

In the traditional model, a business expert explains a need to a product manager, who translates it for a designer, who creates a mockup for an engineer, who builds it. Each handoff introduces delay and translation loss. Critical context disappears. Solutions drift from real needs. IBM's fusion teams invert this: the procurement expert who understands the workflow learns to build on the enterprise AI platform, while the technologist handles integration, APIs, and infrastructure. Both are oriented toward the same business outcome.

IBM also introduced an "AI license to drive" framework—a certification model that allows anyone in the organization to build and deploy AI, regardless of where they sit on the org chart, provided they demonstrate competence. This is a direct structural rejection of the capture dynamic: rather than limiting AI development to the infrastructure department, it democratizes it while maintaining governance through certification rather than gatekeeping. The CIO's own observation is telling: position on the org chart should not dictate who can contribute to AI outcomes.

The Pattern Across the Contrast. The DBS and IBM examples share structural features that map precisely to the governance model this paper advocates. In both cases, business outcomes—not infrastructure metrics—are the primary measure of AI success. Business domain expertise is embedded in AI development teams rather than consulted after the fact. Governance operates as an enabling design principle rather than an external approval gate. And technology functions provide platform infrastructure and

technical standards without claiming ownership of strategy, prioritization, or success metrics.

The contrast with the infrastructure-captured model is stark across every dimension the enterprise research literature tracks. Industry surveys from Deloitte's 2025 State of Enterprise AI report show that organizations where senior leadership actively shapes AI governance—rather than delegating it to technical teams alone—achieve significantly greater business value. McKinsey's 2026 Global Tech Agenda found that nearly two-thirds of top-performing companies have technology leaders deeply involved in crafting enterprise strategy, and these companies are adopting product and platform operating models—structures that embed joint business-technology ownership and cross-functional accountability—at more than four times the rate of others. This is not a speed metric. It is a structural governance metric. The top performers aren't the ones with the largest AI infrastructure budgets or the most data scientists. They are the ones that resolved the strategic governance question correctly—placing business functions in the driver's seat while keeping technology in the engine room.

Importantly, the Deloitte data also reveals a cautionary asymmetry: while organizations that get governance right accelerate, those that get it wrong don't merely stagnate—they accumulate what might be called "governance debt." Shadow AI proliferates. Uncoordinated experiments multiply. Cost structures calcify around platforms and teams that exist to maintain themselves rather than to deliver value. The compounding is real and it runs in both directions: good governance compounds advantage, and bad governance compounds cost. The organizations in the middle—those still debating who owns AI strategy—are falling behind on both curves simultaneously: the advantage curve pulling away above them as leaders compound value, and the cost curve rising beneath them as governance debt accumulates.

The Data: Governance Models and ROI Outcomes

The DBS and IBM cases are instructive, but individual examples—however compelling—are vulnerable to the objection that they represent special circumstances. The broader empirical picture is therefore essential to this paper's argument. Fortunately, the data is now extensive enough to draw directional conclusions, and the direction is consistent: the governance model an organization adopts for AI is a significant predictor of whether AI investments produce measurable business value.

The ROI Landscape: Widespread Disappointment. The baseline context is sobering. MIT's 2025 GenAI Divide report, based on 150 executive interviews, 350 employee surveys, and analysis of 300 public AI deployments, found that approximately 95% of enterprise generative AI pilots failed to produce measurable financial returns. An IBM CEO study found that only about 25% of AI initiatives delivered expected ROI, and just 16% had been scaled enterprise-wide. McKinsey's research shows that 80% of

enterprises report no tangible EBIT impact from their generative AI investments. Deloitte's 2025 survey of 1,854 executives found that only 6% of organizations achieved payback on AI investments within one year, against a typical technology investment expectation of seven to twelve months.

These numbers are striking in their consistency: across different research methodologies, sample sizes, and geographies, the same picture emerges. The majority of enterprise AI investments are not delivering expected financial returns. The standard explanations—immature technology, poor data quality, regulatory uncertainty—are real but insufficient. If these were the primary barriers, we would expect failure rates to be more uniformly distributed across organizational types. They are not.

The Governance Differentiation. When the data is disaggregated by governance maturity, a clear pattern emerges. Organizations with mature, business-aligned governance structures outperform their peers on virtually every measure of AI value creation. An important translation is required here: when the research literature uses terms like "mature," "strong," or "responsible" governance, it is describing organizations that have resolved the ownership question—organizations where strategic and value governance are aligned with business functions, not organizations that have merely built elaborate process governance and compliance governance under IT control. An infrastructure-captured AI program with robust approval gates and thorough compliance reviews may score well on process maturity, but it is not what the research identifies as "mature governance." The maturity that correlates with ROI is structural, not procedural.

With that translation in mind, the data is striking. The IBM Institute for Business Value, surveying 1,000 senior leaders across 17 industries, found that 27% of AI efficiency gains are directly attributable to strong governance, and that companies investing more heavily in AI ethics reported 34% higher operating profit from their AI programs. McKinsey's research found that organizations embedding responsible AI governance from the start see up to 40% higher ROI from AI investments, driven by reduced rework, fewer audit failures, and faster deployment cycles. Competitors with mature governance frameworks deploy AI 40% faster and achieve 30% better ROI than those struggling with governance bottlenecks.

The Cloud Security Alliance's State of AI Security and Governance survey, conducted with Google Cloud, found that governance maturity was the single strongest predictor of real-world AI readiness across all dimensions measured. Organizations with comprehensive governance policies were nearly twice as likely to have adopted advanced AI capabilities compared to those with partial guidelines, and far more likely to have moved AI from experimentation into production. A Gartner survey reinforced the longevity dimension: 45% of organizations with high AI maturity keep AI initiatives

operational for three years or more, compared to just 20% of their lower-maturity peers—and the key differentiator was governance structure, not technology stack.

The Governance Bottleneck Paradox. An important nuance in the data prevents a simplistic reading. The 2025 AI Governance Benchmark Report found that 44% of AI leaders said their governance processes were too slow, and 58% cited disconnected governance systems as the top barrier to scaling AI. This appears to contradict the argument that governance enables value. It does not. It reveals that the dominant form of governance—centralized, IT-controlled, process-heavy—is itself the bottleneck. The organizations reporting governance as too slow are overwhelmingly those operating under the infrastructure-captured model described in this paper: governance owned by IT, structured as an approval gate, and optimized for risk minimization rather than value delivery.

The organizations reporting governance as an accelerator—the DBS Banks and IBMs of the data set—are those that embedded governance within business-led development processes rather than layering it on top as an external checkpoint. The data does not say governance slows AI. It says the wrong kind of governance slows AI. Specifically, it says process governance that is controlled by infrastructure departments, disconnected from value governance, and structured around approval rather than enablement produces the worst outcomes: slower than ungoverned shadow AI, more expensive than lightweight business-led governance, and less effective at managing risk than compliance governance embedded in development processes.

The Compounding Gap. Perhaps the most consequential finding in the current data is the widening gap between organizations that have resolved the governance question and those that have not. Deloitte's research shows that twice as many leaders reported transformative AI impact in 2025 compared to the prior year—but this improvement was concentrated among organizations with mature governance, not distributed across the population. The gap between AI leaders and laggards is widening, not narrowing. Leaders are deploying AI in under three months from concept to production while laggards remain stuck in extended pilot cycles.

This compounding dynamic is the most urgent element of the empirical picture. If the gap were stable, organizations could afford a gradual governance transition. But the data shows acceleration at the top and stagnation in the middle and bottom. The organizations that resolved AI governance early—placing business ownership at the center, embedding governance in development, measuring value rather than activity—are pulling away. The organizations still debating who owns AI strategy, still routing all initiatives through IT-controlled approval processes, still measuring deployment velocity instead of business value, are falling behind at an increasing rate. The

compounding is not linear. It is exponential. And every quarter of unresolved governance is a quarter of widening competitive distance.

The Incentive Problem: Speed, Safety, and the Missing Variable

Understanding why infrastructure capture produces poor AI outcomes requires examining the incentive structures at work. Two dominant imperatives drive most enterprise AI programs today: speed and safety. Executive leadership demands rapid AI adoption to keep pace with competitors and capture productivity gains. Responsible AI and governance teams demand careful risk management to avoid legal, compliance, and reputational exposure. Both imperatives are reasonable in isolation. The problem is that organizations treat speed and safety as proxies for a third variable they are not measuring: business value.

The tension between speed and safety—and the absence of value as a governing variable—has been examined from complementary angles in two related papers. “The Productivity-Risk Tradeoff of Personal AI Prohibition in Knowledge Work” argued that enterprise AI policies built solely on risk avoidance destroy compounded cognitive capital, drive advanced users to shadow AI workarounds, and paradoxically increase the data exposure they are designed to prevent. That paper’s central insight—that over-restriction can increase rather than reduce risk—is the individual-level manifestation of the governance bottleneck and shadow AI dynamics examined in Sections 3.3 and 5.2 of the present paper. Separately, “Slow Down to Speed Up” examined the operational consequences of “go faster with AI” mandates, arguing that organizations should reframe the directive from technology adoption to business outcomes: lift revenue faster, reduce cost faster, minimize risk, and find out quickly whether AI earns its place. The present paper extends both arguments to the structural level: the speed-versus-safety tension and the prohibition-versus-productivity tradeoff are not merely operational challenges to be managed. They are predictable outputs of a governance model in which the wrong function controls the wrong variables. Resolve the structural problem—place strategic and value governance where they belong—and both operational tensions become manageable. Leave the structure misaligned, and no amount of operational optimization will compensate.

The Go-Fast Imperative

The pressure to “go fast with AI” typically originates at the executive level and flows downward through CIO organizations. This pressure creates institutional incentives to demonstrate visible activity: vendor selections, platform deployments, headcount buildouts, pilot launches, and proof-of-concept demonstrations. The metric that matters in this environment is velocity of AI deployment, not quality of business outcomes.

This incentive structure produces predictable distortions. Vendor selections favor platforms that promise broad capability and fast deployment over solutions optimized for specific business outcomes. Professional services engagements balloon as consulting firms supply the "AI expertise" that infrastructure departments lack in business domains. AI teams expand with expensive technical talent whose skills center on model architecture and platform engineering rather than business process understanding. The result is large, expensive AI organizations that can deploy technology rapidly but struggle to deliver measurable business value.

The Safety Imperative

Responsible AI teams serve a legitimate and important function. Legal, compliance, and reputational risks associated with AI are real and require structured governance. The problem is not the existence of these teams but the scope of their authority and the incentive structures under which they operate.

Responsible AI functions, typically embedded within IT governance structures, are incentivized to minimize risk. This is an appropriate objective for compliance governance. It is not an appropriate objective for value governance. A Responsible AI team optimizing for risk minimization will consistently prefer restrictions that reduce exposure at the cost of business utility. They will establish guardrails that prevent harmful outcomes without establishing guardrails that ensure valuable ones. They will create process governance that slows deployment without creating value governance that verifies value delivery.

The result is a governance model that is effective at preventing bad AI outcomes and ineffective at producing good ones.

The Venn Diagram Fallacy

The implicit assumption in most enterprise AI programs is that speed and safety, properly balanced, will produce business value. This assumption is wrong. Speed, safety, and value are three distinct objectives. A Venn diagram of the three shows meaningful overlap in some regions and no overlap in others. An AI deployment can be fast and safe and value-destroying. It can be slow and risky and enormously valuable. The current governance model optimizes for two of the three variables and assumes the third will follow. It doesn't.

What's missing is value governance—a function accountable for business outcomes not as a secondary benefit of fast, safe deployment, but as the primary objective against which both speed and safety are calibrated. This function does not exist in most enterprise AI organizations because the departments that would naturally own it—the

business functions closest to the processes being transformed—have been subordinated to the infrastructure departments driving the AI program.

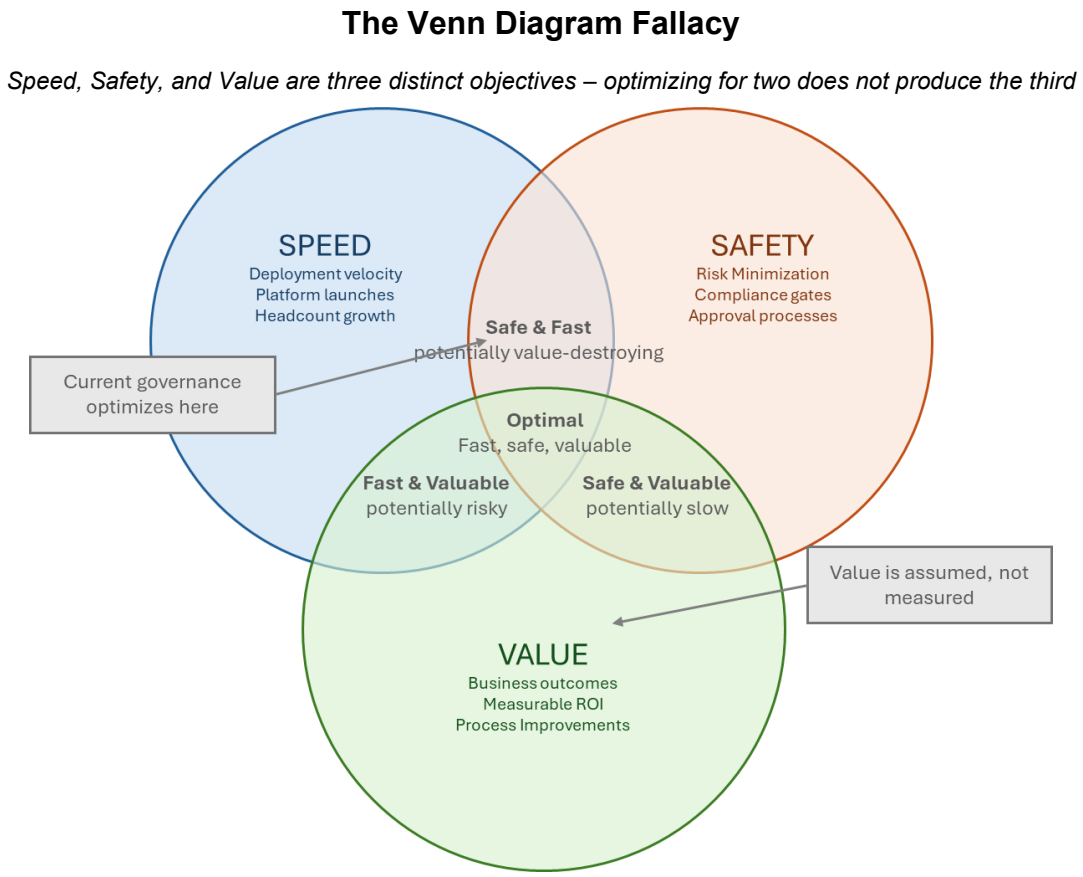

Figure 2. The Venn Diagram Fallacy — The current governance model treats Speed and Safety as proxies for Value. They are not.

The Cost Structure Problem

The financial consequences of infrastructure capture are visible in the cost structures of enterprise AI programs. Organizations that have allowed IT departments to drive AI strategy consistently exhibit a characteristic pattern of cost inflation that depresses ROI.

First, headcount bloat. AI teams within IT organizations grow rapidly with expensive technical specialists—data scientists, ML engineers, platform architects—whose skills are oriented toward building and maintaining AI infrastructure rather than understanding business processes. These teams are often staffed to technical-complexity assumptions that reflect IT's perspective on what AI requires, rather than business-outcome assessments of what AI needs to deliver.

Second, professional services dependency. Infrastructure-oriented AI teams frequently lack the business domain expertise to design effective AI solutions, leading to heavy

reliance on external consulting and systems integration firms. These engagements are expensive, extend timelines, and often produce solutions optimized for the consulting firm's platform preferences rather than the organization's business needs.

Third, platform over-engineering. IT-driven AI programs tend to invest heavily in enterprise AI platforms, data infrastructure, and tooling that provide broad technical capability but are poorly matched to specific business value opportunities. The platform becomes the strategy, rather than the strategy driving platform decisions.

Fourth, and perhaps most consequentially, development velocity mismatch. This cost driver deserves extended treatment because it represents not merely one more item on a list of infrastructure capture consequences, but a structural transformation in what AI development requires—one that fundamentally undermines the operating model of IT-governed AI programs.

The Velocity Mismatch in Detail

Traditional enterprise IT development operates on a cadence shaped by decades of software engineering practice: requirements gathering, architecture review, design, development, testing, security assessment, staged rollout. This process evolved for good reasons—it manages complexity, reduces risk, and ensures quality in systems that are expensive to build and difficult to modify once deployed. For ERP implementations, infrastructure migrations, and large-scale systems integration, this cadence remains appropriate.

AI development operates on a fundamentally different cadence. The availability of pre-trained foundation models, commercial AI APIs, low-code and no-code AI platforms, and modern prompt engineering techniques has compressed the build cycle for many AI applications from months to weeks—and in some cases from weeks to days. In 2026, custom AI solutions routinely move from concept to production in four to eight weeks. Proofs of concept that once required dedicated data science teams and months of model training can now be assembled by a business analyst with domain expertise and access to commercially available tools. Google's CEO disclosed that AI now generates over 25% of the company's new code. AI-assisted development has reduced coding timelines by more than half according to multiple industry analyses. The technology has changed what building with AI looks like.

This is not an incremental acceleration. It is a category change in what AI development requires, and it has three implications that infrastructure-captured AI programs are structurally unable to absorb.

First, the build calculus has inverted. When AI development required months of specialized work by expensive technical teams, centralizing that work within IT made

economic sense—you concentrated scarce expertise where it could be managed efficiently. When AI development requires weeks of work by domain experts using commercial tools, centralization becomes a bottleneck rather than an efficiency. The scarce resource is no longer AI engineering talent. It is business domain expertise—the knowledge of how specific processes work, what outcomes matter, and where value hides in operational workflows. This expertise already lives in business functions. Routing AI development through an intermediary department that lacks this expertise, only to have that department hire consultants to supply it, is an expensive detour that produces worse outcomes on a longer timeline.

Second, the iteration cycle has compressed below the governance cycle. When AI development took six to twelve months, an eleven-week process governance review represented a manageable proportion of the total timeline—painful but not fatal. When AI development takes four to eight weeks, an eleven-week process governance review is longer than the development itself. The process governance cycle is no longer a gate within the development cycle. It is a wall that the development cycle crashes into. Business teams that can prototype a working AI solution in three weeks are told to wait three months for approval to deploy it. By the time the approval arrives, the underlying models have improved, the competitive landscape has shifted, and the solution needs to be rebuilt with current capabilities rather than deployed in its now-outdated form. The composite example of the marketing documentation team in Section 3.3 illustrates this precisely: by the time the IT-governed process delivered an approved solution, it was solving yesterday's problem with yesterday's technology.

It is worth being concrete about what a "business team" capable of this velocity actually looks like. It is not a skunkworks of rogue developers. It is a small, cross-functional unit: a business lead with solid understanding of enterprise objectives and strategy, a business analyst or subject matter expert with deep domain knowledge of the process being transformed, an AI-capable resource who has been certified on the enterprise's AI platform (the "license to drive" concept from IBM), and direct access to IT integration resources who handle the technical plumbing—API connections, data lake access, platform integrations. These roles can overlap in individuals; a business SME with Salesforce Einstein Prompt Builder expertise, for instance, can move quality automation into production far more quickly than the dedicated development teams typically assembled for AI projects under IT governance. That is the entire purpose of low-code and no-code AI tools—development is democratized so that the people closest to the business process can build solutions without waiting in an IT development queue. In infrastructure-captured environments, IT often constrains these tools to operate far below their designed capabilities, requiring developer intermediation for workflows the platform was built to handle without one.

Third, AI development is increasingly self-accelerating. AI tools are being used to build AI solutions. Code generation, automated testing, deployment automation, and monitoring are all being transformed by the same technology they support. This creates a compounding velocity advantage for teams that operate within AI-native development frameworks: each generation of tools makes the next generation faster to build. Teams operating within traditional IT governance structures miss this compounding because their development cadence is gated by process timelines, not technical capability. The infrastructure department's sprint cycle does not accelerate because AI can now write code faster. The architecture review board does not meet more frequently because prototypes can be built in days. The security assessment does not complete sooner because the underlying tools have improved. The governance infrastructure is temporally fixed while the development capability is temporally accelerating. The gap between what could be built and what is permitted to be built widens with every quarter.

The Shadow AI Consequence

The predictable consequence of this velocity mismatch is the same consequence that has followed every previous technology governance mismatch in enterprise history: shadow adoption. Business teams that are told to wait months for IT-governed AI deployment do what business teams have always done when infrastructure departments cannot keep pace with business needs—they route around the obstruction. Commercial AI tools are inexpensive, accessible, and require no IT involvement to use. The result is widespread unsanctioned AI usage that is invisible to governance structures, unmeasured by the organization, and unmanaged from a risk perspective.

Industry data confirms this pattern at scale. Surveys indicate that more than three-quarters of knowledge workers now use AI tools, and a similar proportion bring their own AI tools to work without employer oversight. The amount of corporate data flowing into AI tools has increased dramatically year over year. This is not an education problem or a compliance problem. It is a governance design problem. When the sanctioned path to AI deployment takes months and the unsanctioned path takes minutes, rational actors will choose the unsanctioned path. Shadow AI is not a cause of governance failure. It is a symptom of velocity mismatch—and velocity mismatch is a structural consequence of infrastructure capture.

The Build-vs.-Buy Inversion

The velocity transformation has also inverted the traditional enterprise build-versus-buy decision in ways that infrastructure-captured AI programs are poorly equipped to navigate. The historical IT model favored building custom solutions when competitive differentiation was required and buying when commodity functionality sufficed. AI has scrambled this logic. The majority of enterprises now purchase or subscribe to AI solutions rather than building internally, because commercial AI platforms deliver

capability faster, at lower cost, and with less technical debt than custom builds. Industry data consistently shows that organizations combining vendor tools with internal business domain expertise scale AI significantly faster than those attempting fully custom development.

Infrastructure-captured AI programs resist this shift because custom building justifies the large technical teams, platform investments, and professional services engagements that constitute the AI program's cost structure. An AI team of ninety people maintaining a custom platform has an institutional incentive to build rather than buy, regardless of which approach delivers value faster. The buy decision threatens the team's headcount. The build decision sustains it. This is not a conspiracy—it is a predictable consequence of incentive structures. When the department that selects AI approaches is the same department whose budget and headcount depend on those approaches being complex, the selection will consistently favor complexity over simplicity, custom over commercial, and platform over purpose.

These cost structures are not incidental—they are structural consequences of infrastructure capture. When the department controlling AI strategy is incentivized to build and maintain complex technical infrastructure, the AI program will be expensive, infrastructure-heavy, and slow. The business value proposition is buried under layers of technical architecture that exist to serve IT's operating model, not the business's value creation needs.

Toward a Resolution: Business-Owned AI Strategy

The web ownership wars ended when organizations recognized that the departments closest to customers and business outcomes needed to own the strategy for customer-facing technology. The same resolution is needed for AI, extended to encompass internal business processes as well.

The Target Governance Model

The governance model this paper advocates is structurally simple: business units own AI strategy, prioritization, and outcomes. Technology functions provide enabling infrastructure, platform services, and technical governance (security, compliance, architecture standards). The critical distinction is between enabling and driving. IT enables AI deployments by providing reliable infrastructure, enforcing security requirements, and ensuring architectural coherence. Business functions drive AI deployments by identifying value opportunities, defining success metrics, selecting approaches optimized for business outcomes, and being accountable for ROI.

This model does not diminish IT's importance. Infrastructure, security, and technical governance are essential functions. But they are enabling functions, not strategic ones.

The web taught us that infrastructure departments lack the customer proximity and business domain expertise to make good strategic decisions about technology that serves customers and business processes. AI is teaching us the same lesson at greater scale and higher cost.

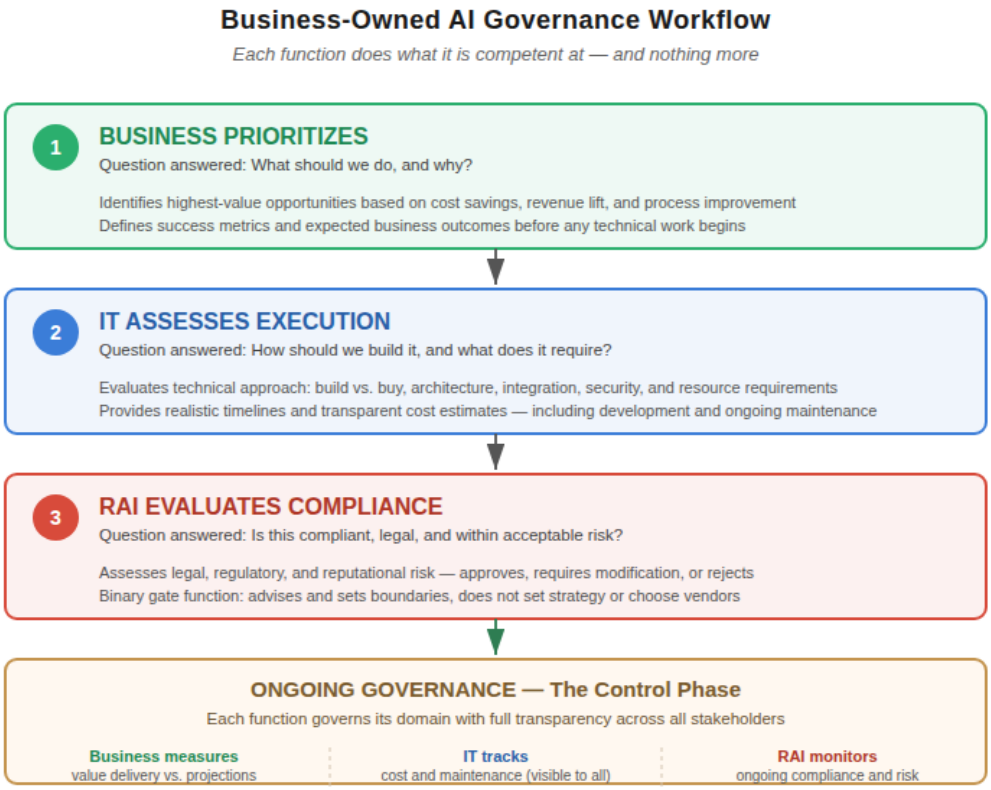

Figure 3. Business-Owned AI Governance Workflow — Each function does what it is competent at, and nothing more.

The element of this workflow most absent from current practice is the Control phase at the bottom—what this paper has defined as value governance. To illustrate what value governance looks like in practice, consider a concrete example. A business unit projects that an AI-driven automation of its invoice processing workflow will save $600,000 annually in labor costs and reduce processing errors by 40%. Under current governance, the initiative would be approved based on this projection, deployed by IT, assessed for compliance by RAI, and then—in most organizations—never rigorously measured against the original business case.

Under value governance, the story continues past deployment. Six months in, the business measures actual labor cost reduction and error rates against projections. But critically, IT also reports the fully loaded cost of the deployment: $350,000 in development, $90,000 in annual platform licensing, and $75,000 in ongoing maintenance and support. These numbers are visible to all stakeholders—not buried in

IT's operating budget where the business unit never sees them. The net calculation is straightforward: $600,000 in projected savings minus $165,000 in annual IT costs yields $435,000 in net value—if the savings projections hold. If actual savings come in at $300,000, the initiative is barely breaking even on a fully loaded basis, despite appearing successful when only the business-side metrics are visible.

This is the discipline that value governance introduces: the same financial rigor that any other capital investment would receive, applied to AI deployments with full cost transparency across all stakeholders. It is not a new concept—it is standard practice in manufacturing, capital budgeting, and operational management. Its absence from enterprise AI programs is not an oversight. It is a structural consequence of infrastructure capture: when the department incurring the costs is the same department controlling the governance process, those costs remain opaque to the business functions accountable for value delivery. Making IT costs transparent is not punitive toward IT. It is the basic financial hygiene that allows an organization to know whether its AI investments are actually working.

Responsible AI Under Business Ownership

A common objection to business-owned AI strategy is that business functions will prioritize speed and value over safety. This concern is legitimate but misframed. The current model—in which Responsible AI operates as an extension of IT governance—is effective at managing legal and compliance risk but ineffective at aligning AI deployment with business value. A better model would maintain independent Responsible AI oversight for legal, compliance, and reputational risk while embedding value-oriented governance within business ownership.

Responsible AI should function like a corporate legal team: it advises, it sets boundaries on genuinely risky behavior, and it ensures regulatory compliance. It does not set business strategy, choose vendors, or determine which processes to transform. When Responsible AI operates within an IT governance structure, it tends to absorb the host department's risk-minimization incentives and apply them to domains where value creation, not risk avoidance, should be the governing objective.

The Urgency of Intentional Resolution

The web ownership question resolved itself eventually, through organic organizational learning. This paper argues that waiting for organic resolution with AI is a mistake organizations cannot afford. The costs of misaligned governance compound more rapidly with AI than they did with the web, because AI touches more organizational surface area and because the competitive landscape is moving faster.

Organizations that establish business-owned AI governance now will compound the advantages of aligned incentives and business-proximate decision-making. Organizations that allow infrastructure capture to persist will compound the costs of misaligned governance, bloated cost structures, and value-destroying deployment patterns. The gap between these two trajectories widens with every quarter.

Addressing the Counterarguments

The governance model proposed in this paper invites several reasonable objections. Taking them seriously—rather than dismissing them—is essential both to intellectual honesty and to practical implementation. Three objections deserve direct engagement.

Objection 1: Business units lack the technical judgment to own AI strategy. This is the most common and most intuitive objection. AI is technically complex. Business teams do not typically have deep expertise in model architecture, data engineering, security, or systems integration. Entrusting AI strategy to people who lack technical fluency risks poor decisions, wasted investment, and unmanaged risk.

The objection is reasonable on its face but rests on a conflation that this paper exists to disentangle: the conflation of strategic ownership with technical execution. Business ownership of AI strategy does not mean business teams build AI systems without technical support. It means business teams define what problems to solve, how to measure success, which approaches best serve business outcomes, and whether the results justify the investment. Technology teams execute the technical work—building, integrating, securing, and maintaining AI systems—within the strategic direction set by the business.

This is not a novel organizational concept. It is how well-run companies already operate in nearly every other domain. Marketing sets brand strategy; agencies and production teams execute. Product management defines the roadmap; engineering teams build it. Finance sets capital allocation priorities; operations teams deploy resources accordingly. In none of these cases does the executing function own the strategy, and in none of these cases does anyone argue that the strategic function needs to understand production techniques at an expert level. The objection that business units lack technical judgment applies only if we accept the premise that AI strategy is a technical discipline. This paper argues it is a business discipline with technical requirements—a critical distinction.

Moreover, the empirical record suggests that the objection's premise is backwards. The cases examined in this paper—DBS Bank's two-in-a-box model, IBM's fusion teams—demonstrate that placing business and technology in shared governance produces better outcomes than either function operating alone. DBS did not hand AI strategy to bankers with no technical support. It paired business leaders with technology leaders, gave them shared KPIs, and held both accountable for business outcomes. IBM did not tell

procurement specialists to build AI systems without help. It paired them with CIO technologists who handled integration and infrastructure while the domain experts shaped the solution. In both cases, business ownership did not exclude IT. It included IT in a governance structure oriented toward business value.

This reveals the deeper asymmetry the objection conceals. Business-owned AI governance includes technology functions at the table—as essential partners in execution, technical governance, and infrastructure. IT-owned AI governance tends to exclude business functions from strategic decisions—or consults them superficially after strategic choices have already been made. The question is not whether IT has a role. The question is whether IT's role is to drive strategy or to enable it. Business-owned governance builds a bridge between business domain expertise and technical capability. IT-owned governance builds a moat that protects infrastructure department interests while restricting the access of the business functions closest to the outcomes that AI is meant to produce.

There is a more fundamental reason this objection is losing force with each passing quarter: the technical skill floor for AI development is dropping rapidly. When prompt-writing skills become as ubiquitous as the ability to build an Excel pivot table—and we are approaching that threshold now—the argument that only IT can build with AI becomes untenable. In this environment, IT's essential technical contribution narrows to what it has always done best: integrating APIs, managing data lake access, securing platform connections, and maintaining the infrastructure that business-built solutions run on. A business analyst building automation in Salesforce Einstein Prompt Builder does not need a data engineer to write the logic. She needs IT to ensure the platform connects securely to the right data sources. That is a genuinely important enabling function. It is not a strategic governance function. Conflating the two is how infrastructure capture persists even as the technical rationale for it evaporates.

Objection 2: This paper is anti-IT and anti-CIO. It is not. This paper argues that CIOs and IT leaders are behaving rationally within incentive structures that reward infrastructure expansion and risk minimization. The problem is structural, not personal. CIOs who capture AI strategy are doing what their institutional position encourages them to do. Criticizing individuals for responding to incentives is less productive than redesigning the incentives themselves.

The most effective CIOs in the current landscape are the ones who have recognized this dynamic and proactively restructured their organizations' relationship with AI—moving from ownership to enablement, from gatekeeping to platform provision, from strategy authorship to strategy execution. IBM's CIO explicitly articulated this shift: where you sit on the org chart should not dictate what you can contribute. The CIOs who will thrive in the emerging governance model are those who embrace the enabling role—not

because it diminishes their importance, but because it positions technology leadership as the essential infrastructure that makes business-owned AI strategy possible. The CIO who provides a reliable, secure, well-governed AI platform and empowers the entire organization to build value on it is more strategically important than the CIO who controls a centralized AI team that builds solutions the business didn't ask for.

Objection 3: The technology is still immature, and these governance questions are premature. This objection holds that AI is evolving so rapidly that establishing governance structures now risks locking in models that will be obsolete within months. Better to wait for the technology to stabilize before investing in governance design.

The web parallel directly refutes this argument. Organizations that waited for the web to "stabilize" before resolving ownership questions lost years of competitive advantage to those that established governance models early—models that evolved with the technology rather than waiting for the technology to stop evolving. The web never stabilized in any meaningful sense. It continuously transformed, and organizations with clear ownership structures adapted to each transformation faster than those still debating who was in charge.

AI will not stabilize either. Model capabilities, development frameworks, deployment patterns, and regulatory requirements will continue to evolve rapidly. This is precisely why governance structures must be established now. The governance model this paper proposes—business ownership of strategy, IT enablement of execution—is not a specific technical architecture that might become obsolete. It is an organizational design principle that remains valid regardless of which specific AI technologies prevail. Business proximity to outcomes, shared accountability for value, and clear distinction between strategic and enabling functions are governance principles that transcend any particular technology generation. Waiting for technological stability to resolve organizational governance is like waiting for the weather to stabilize before building a roof. The instability is the reason the structure is needed, not a reason to defer it.

Conclusion

Enterprise AI's ROI problem is not primarily a technology problem. It is a governance problem with a well-documented historical precedent. Infrastructure departments captured web strategy in the late 1990s and early 2000s, extending legitimate infrastructure authority into strategic domains where they lacked customer proximity and business expertise. The same capture is occurring with AI, amplified by AI's deeper penetration into internal business processes.

The resolution is the same resolution that emerged from the web wars: business ownership of strategy and outcomes, with technology functions serving as enabling

infrastructure. The difference is that organizations cannot afford to reach this resolution through years of organic correction. The competitive costs of misaligned AI governance are too high, the pace of technological change is too fast, and the compounding effects of good and bad governance are too powerful.

The CIOs and IT leaders executing this capture are not villains. They are behaving rationally within incentive structures that reward infrastructure expansion and risk minimization. The problem is structural, and the solution must be structural: intentional governance design that places AI strategy ownership where it belongs—with the business functions accountable for the outcomes AI is meant to produce.

About the Author

Pat Brenner is a senior leader in digital transformation, AI-enabled operations, and enterprise customer engagement. Pat has extensive experience guiding regulated organizations through technology adoption at scale, with a focus on balancing productivity, risk, and human-centered design. His research portfolio on SSRN spans AI governance, deployment strategy, and cultural analysis. Contact: patb@earthoffice.net

Sources and References

Companion Papers

Brenner, P. “The Productivity–Risk Tradeoff of Personal AI Prohibition in Knowledge Work.” Position paper, 2025. Published on LinkedIn and earthoffice.net; SSRN submission forthcoming.

Section: AI Replay, Composites, Comparative Analysis, and Empirical Data

Challapally, A. et al. “The GenAI Divide: State of AI in Business 2025.” MIT Project NANDA, 2025. Basis for the 95% pilot failure rate cited in Section 3.5.

IBM Institute for Business Value. “How Governance Increases Velocity.” IBM, 2025. Survey of 1,000 senior leaders; source for the 27% efficiency gains and 34% higher operating profit findings in Section 3.5.

IBM CEO Study, 2025. Source for the finding that 25% of AI initiatives deliver expected ROI and 16% scale enterprise-wide.

Deloitte AI Institute. “State of AI in the Enterprise, 2026.” Survey of 3,235 leaders, August–September 2025. Source for governance-maturity and ROI correlation findings in Section 3.5.

Deloitte. “AI ROI: The Paradox of Rising Investment and Elusive Returns.” Survey of 1,854 executives, 2025. Source for the 6% payback-within-one-year finding.

Cloud Security Alliance and Google Cloud. “The State of AI Security and Governance.” December 2025. Source for governance maturity as strongest predictor of AI readiness.

McKinsey & Company. “The New CIO Mandate: Strategy, Speed, and Scaled Intelligence” (Global Tech Agenda 2026). February 2026. Source for top-performer governance models and product/platform adoption rates.

ModelOp. “2025 AI Governance Benchmark Report.” Survey of 100 senior AI and data leaders. Source for the 44% “governance too slow” and 58% “disconnected systems” findings.

Gartner, Inc. AI maturity survey, 2025. Source for the finding that 45% of high-maturity organizations maintain AI initiatives for 3+ years vs. 20% of lower-maturity peers.

Section: DBS Bank and IBM Case Studies

McKinsey & Company. “An Inside Look at How McKinsey Helped DBS Become an AI-Powered Bank.” Interview with Piyush Gupta, CEO, DBS Bank. Source for the two-in-a-box model, 33 platforms, and shared KPI structure.

DBS Bank. “DBS’ AI-Powered Digital Transformation.” dbs.com. Source for SGD 180M (2022) economic value, 600+ AI/ML models, and PURE governance framework.

Tan, S.S. (CEO, DBS Bank). Remarks at Singapore Fintech Festival, November 2025. Reported by CNBC. Source for SGD 1B (2025) AI revenue contribution and 1,500+ models in production.

Zhu, F., Zhu, H., and Wong, A. “DBS’ AI Journey.” Harvard Business School Case 625-053, August 2024.

Lyteson, M. (CIO, Technology Platform Transformation, IBM). “Scaling Enterprise AI: Lessons in Governance and Operating Models from IBM.” Stack Overflow Leaders of Code podcast, January 2026. Source for AI fusion teams, AI license to drive, and procurement team examples.

Section: Development Velocity and Shadow AI

Pichai, S. (CEO, Alphabet/Google). Disclosure that AI generates over 25% of Google’s new code. Widely reported, November 2025.

Microsoft and LinkedIn. “2024 Work Trend Index.” Source for the finding that 75% of knowledge workers use generative AI and 78% bring their own AI tools without employer oversight.

PwC. “2025 AI Business Predictions.” Source for Responsible AI survey findings: 60% cite ROI and efficiency benefits, 55% report improved customer experience.